
Project 1.1 Implementation

Amit Menon

Department of Computer Science
State University of New York at Buffalo, NY
amitanil@buffalo.edu

Abstract

The aim is to compare the two approaches (Software 1.1 and Software 2.0) to solve the FizzBuzz problem and analyze the performance of the Software 2.0 solution in effect to changes in certain parameters.

Overview

To solve the Fizz Buzz problem a multiclass classification model with stochastic gradient descent optimization is used. We have a training data of numbers from 101 to 1000 labelled Fizz, Buzz, FizzBuzz and other as specified. We process the input into their right shifted bit representation and take each bit as a feature for input to train the model and process the output labels as 4 numbers representing 4 classes in the problem. This input is then fed to the neural network which has only one hidden layer with 100 nodes initially and uses ReLu as activation function for each node. This model is then optimized using gradient descent optimizer for which the learning rate is set to 0.5. To make computation faster we train the model with batches of training data selected randomly and iterate over all the training data for 'epoch' number of times. On each iteration we thus compare the output layer of the neural network with the given label of the input and make adjustments to reduce the error, thus improving upon each iteration and finally reaching the lowest error margin

1 Modification of Hyperparameters

1.1 Modification of Epoch and Batch Sizes

Epoch is one complete cycle of the whole training data. Beginning with a lower value of 4000 we find that accuracy is not able to peak in 4000 iterations. That implies that the model was not able to train itself efficiently in said cycles. This is called underfitting of the model as shown in the below graph.

Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x1e88471a320>




Figure 1.1 a)

As seen in the graph, the accuracy falters and rises quickly, and doesn't max out completely reaffirming the observation, that the model isn't well trained. Thus on increasing the epoch to 8000 we get Fig. 1.1. b). On this the graph peaks out in the region of 5000 and then radically thins down. This implies that at this point the model is trained to the extent that it won't work on any data apart from training data. This is overfitting. Also its clear that that the optimum value lies in the region of 5000 iterations so taking epoch as 5500 we obtain Fig 1.1. c) and take epoch = 5500 as fixed value for modification of other parameters.




Figure 1.1 b)

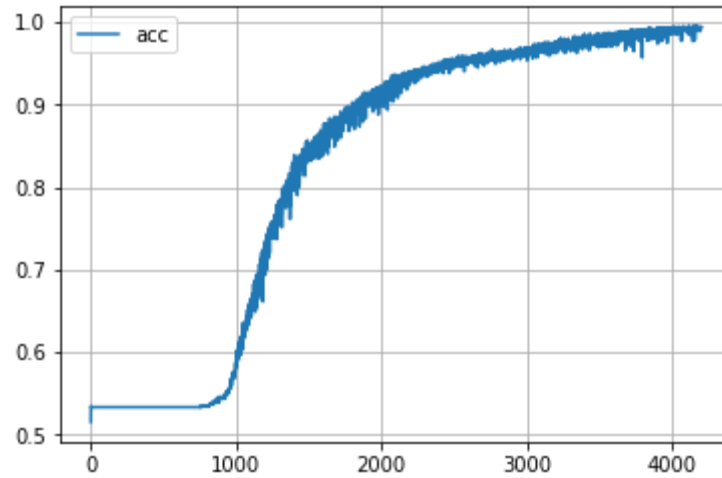


Figure 1.1 c)

Table 1: Accuracy results for variations on epoch

EPOCH	ACCURACY
5500	95
5500	95
5500	97
8000	94
8000	99
8000	96
4000	91
4000	98

Batches are the subset of data for which one sets of calculations are made for ease of calculation. Tweaking this value first to 90 gave Fig 1.1.d). The thickness of this plot implies that accuracy peaked and dropped quite a few times which implies that it was not optimally trained. So we increase the batch size to 150 to obtain Figure 1.1.e). This again reflects a thick plotting of the accuracy graph and doesn't peak out, which implies that the batch size is greater than optimal. So we reduce the size again to 90 which gives result as Fig 1.1.f). Since this plots the thinnest curve, and converges correctly in the 5000 region itself, batch size of 100 is chosen as the optimum value.




Figure 1.1 d) Batch size 90




Figure 1.1 e) Batch size 150




Figure 1.1 f) Batch size 100

Table 2: Comparison of Batch and Accuracy

BATCH	ACCURACY	EPOCH
-------	----------	-------

100	96	5500
150	80	5500
200	91	5500
90	92	5500

1.2 Modification of values of Learning rate.

We are using a stochastic gradient descent to correct our model and in that regards alpha is a factor that determines how quickly or slowly will our gradient descent converge to the minimum.

On initially tweaking the learning rate to 0.01, it is found that gradient descent took few number of iterations to even began converging to the local minimum as shown in the graph of accuracy vs iterations below:




Figure 1.2. a) Alpha 0.01

This is an indication that the learning rate is too slow. If the learning rate is too small, gradient descent will not converge in given number of iterations.




Figure 1.2. b) Alpha 0.08

On gradually increasing learning rate, gradient descent began to converge in less number of iterations until it takes extremely big jumps to converge and may skip the local minima altogether as shown in Fig 1.2 b). As we can see that in the plot the climb is way too steep which means gradient descent is taking the steps towards minima quickly in the right

86 directions (less number of iterations).

87

88 1.3 Number of Neurons nodes in Hidden Layer

89 Number of nodes in the hidden layer is a very impactful factor in improving the model. As
90 each node of a layer is an input to node in the next layer. As seen below in Fig 1.3.a) an
91 increase in the number of nodes in the hidden layer makes the model reach greater accuracy
92 in around 2000 iterations itself and without excessively thickening in between iterations.
93 (Jumps are gradual in gradient descent).

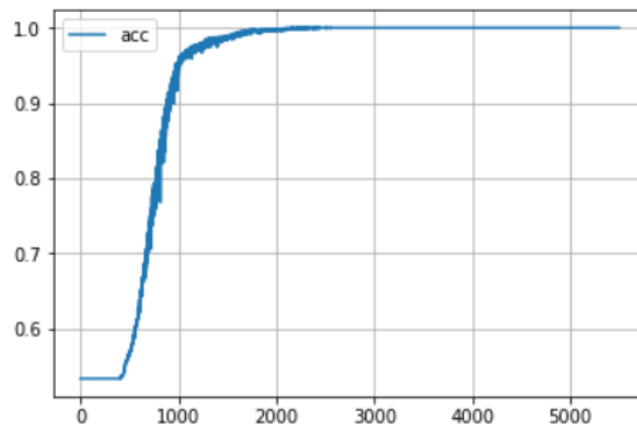
94 The reason for this performance is because increase in number of nodes greatly increases the
95 number of times cost computation is made at each layer and also allows gradients to back
96 propagate in the network to influence values of the node in previous layer towards the value
97 of corresponding node in output layer.

98

No of nodes	ACCURACY	EPOCH	BATCH
128	91	5500	100
128	92	5500	100
256	100	5500	100
256	99.1	5500	100
256	97.3	5500	100
256	97	5500	100

99

```
Out[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1e7589997b8>
```



100
101

Figure 1.3. a)

102 3 Modification of Activation Functions

103 We now analyze the performance impact when we change the activation function for the
104 given hidden layer.

105

106 3.1 Leaky_relu

107 The problem in using ReLu is that sometimes it can cause a node to die out (Dying relu
108 problem). This happens when a large gradient back propagating can cause a node to be
109 updated in such a way that it won't fire at input point for any layer again. To solve this
110 Leaky ReLu is used. Employing leaky relu as an activation function greatly improved

111 performance of the model and the fluctuation in accuracy for different runs as seen below.


112

113 *Table 3:Effect on Accuracy using Leaky Relu*

No of nodes (Leaky Relu)	ACCURACY	EPOCH	BATCH	ALPHA
256	98	5500	100	0.2
256	98	5000	100	0.2
256	99.1	5500	100	0.2
256	93	5500	100	0.2
256	98	5000	100	0.5
256	100	5000	100	0.5

114

115 It is worth noting that tweaking the parameter “alpha” for leaky relu, further improves the
116 accuracy though it is not completely reliant. Also increase in this alpha greatly increases
117 computation time as not all of the nodes are fired. However if the value of alpha is increased even
118 further then it results in over fitting as shown in figure 2 c) for which alpha = .75



119

120

Figure 2 a) Leaky Relu

Out[115]: <matplotlib.axes._subplots.AxesSubplot at 0x1e888589e10>




Figure 2 b) Leaky Relu overfitting

4 Adding a dropout

To prevent overfitting, we use the concept of dropout, which is essentially dropping outputs from nodes from hidden layer on the basis of certain probability which here is 0.99. Notice that with dropouts even as accuracy peaks, it doesn't thin out completely which means that use of dropouts is allowing the model to generalize

Out[147]: <matplotlib.axes._subplots.AxesSubplot at 0x1e88aeec390>




Figure 4 Accuracy with dropout

5 Conclusion

In conclusion, the optimum solution to the Fizz Buzz problem depends upon all the selected parameters and to find a balance in tweaking them. Raising even one of these parameters to a slight extreme can impact the performance and accuracy of the model, as seen in case of over fitting and under fitting of data. The solution also depends on the algorithm chosen for optimization and neural network is implemented (with or without dropout, activation function).

139 Also accuracy here is not a complete indication of how well the model is trained since it runs on
140 randomized input and will give different results on each run, we aim for a consistency over a
141 number of runs to determine how well the model is trained. In our tests we found that the model
142 gave consistent results for the following configuration.

143

No of nodes (Leaky Relu) without dropout	ACCURACY	EPOCH	BATCH	ALPHA	Learning rate
256	98	5500	100	0.2	0.08
256	98	5000	100	0.2	0.08
256	99.1	5500	100	0.2	0.08
256	93	5500	100	0.2	0.08
256	98	5000	100	0.5	0.08
256	100	5000	100	0.5	0.08

144

145